

```
In [1]: # Import tensorflow
import tensorflow as tf
```

```
In [2]: # Load the IMDB dataset
imdb = tf.keras.datasets.imdb
```

```
In [3]: """The argument num_words=10000 means you'll only keep the top 10,000
most frequently occurring words in the training data. Rare words will be discarded. This allows
you to work with vector data of manageable size.
The variables train_data and test_data are lists of reviews; each review is a list of
word indices (encoding a sequence of words). train_labels and test_labels are
lists of 0s and 1s, where 0 stands for negative and 1 stands for positive:"""

(train_data, train_labels), (test_data, test_labels) = imdb.load_data(
    num_words=10000)
```

In [4]: `train_data[0]`

```
Out[4]: [1,  
        14,  
        22,  
        16,  
        43,  
        530,  
        973,  
        1622,  
        1385,  
        65,  
        458,  
        4468,  
        66,  
        3941,  
        4,  
        173,  
        36,  
        256,  
        5,  
        25,  
        100,  
        43,  
        838,  
        112,  
        50,  
        670,  
        2,  
        9,  
        35,  
        480,  
        284,  
        5,  
        150,  
        4,  
        172,  
        112,  
        167,  
        2,  
        336,  
        385,  
        39,  
        4,  
        172,  
        4536,  
        1111,  
        17,  
        546,  
        38,  
        13,  
        447,  
        4,  
        192,  
        50,  
        16,  
        6,  
        147,  
        2025,
```

19,
14,
22,
4,
1920,
4613,
469,
4,
22,
71,
87,
12,
16,
43,
530,
38,
76,
15,
13,
1247,
4,
22,
17,
515,
17,
12,
16,
626,
18,
2,
5,
62,
386,
12,
8,
316,
8,
106,
5,
4,
2223,
5244,
16,
480,
66,
3785,
33,
4,
130,
12,
16,
38,
619,
5,
25,
124,
51,

36,
135,
48,
25,
1415,
33,
6,
22,
12,
215,
28,
77,
52,
5,
14,
407,
16,
82,
2,
8,
4,
107,
117,
5952,
15,
256,
4,
2,
7,
3766,
5,
723,
36,
71,
43,
530,
476,
26,
400,
317,
46,
7,
4,
2,
1029,
13,
104,
88,
4,
381,
15,
297,
98,
32,
2071,
56,
26,

```
141,  
6,  
194,  
7486,  
18,  
4,  
226,  
22,  
21,  
134,  
476,  
26,  
480,  
5,  
144,  
30,  
5535,  
18,  
51,  
36,  
28,  
224,  
92,  
25,  
104,  
4,  
226,  
65,  
16,  
38,  
1334,  
88,  
12,  
16,  
283,  
5,  
16,  
4472,  
113,  
103,  
32,  
15,  
16,  
5345,  
19,  
178,  
32]
```

```
In [5]: train_labels[0]
```

```
Out[5]: 1
```

```
In [6]: """You can't feed lists of integers into a neural network. You have to
turn your lists into
tensors. One-hot encode your lists to turn them into vectors of 0s and
1s. This would
mean, for instance, turning the sequence [3, 5] into a 10,000-dimensional vec-
tor that would be all 0s except for indices 3 and 5, which would be 1
s. Then you
could use as the first layer in your network a Dense layer, capable of
handling
floating-point vector data."""

# Encoding the integer sequences into a binary matrix
import numpy as np

def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1.
    return results

x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
y_train = np.asarray(train_labels).astype('float32')
y_test = np.asarray(test_labels).astype('float32')
```

```
In [7]: # Here's what the samples look like now:
x_train[0]
```

```
Out[7]: array([0., 1., 1., ..., 0., 0., 0.])
```

In [8]: *"""There is more to this but for the time being, you'll have to trust me with the following architecture choice:."""*

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(16, activation='relu', input_shape=(10000,)),
    tf.keras.layers.Dense(16, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')])
```

"""Two intermediate layers with 16 hidden units each. And a third layer that will output the scalar prediction regarding the sentiment of the current review"""

"""The intermediate layers will use relu as their activation function, and the final layer will use a sigmoid activation so as to output a probability (a score between 0 and 1, indicating how likely the sample is to have the target "1": how likely the review is to be positive). A relu (rectified linear unit) is a function meant to zero out negative values."""

Out[8]: 'The intermediate layers will use relu as their activation function, and the final layer\nwill use a sigmoid activation so as to output a probability (a score between 0 and 1,\nindicating how likely the sample is to have the target "1": how likely the review is to be\npositive). A relu (rectified linear unit) is a function meant to zero out negative values.'

In [9]: *# Compiling the model*

```
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['accuracy'])
```



```
In [10]: # Setting aside a validation set
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]

history = model.fit(partial_x_train, partial_y_train, epochs=20, batch
_size=512, validation_data=(x_val, y_val))
```

Train on 15000 samples, validate on 10000 samples

Epoch 1/20

15000/15000 [=====] - 1s 95us/sample - loss: 0.5086 - accuracy: 0.7908 - val_loss: 0.3907 - val_accuracy: 0.8665

Epoch 2/20

15000/15000 [=====] - 1s 45us/sample - loss: 0.3096 - accuracy: 0.9007 - val_loss: 0.3097 - val_accuracy: 0.8838

Epoch 3/20

15000/15000 [=====] - 1s 46us/sample - loss: 0.2212 - accuracy: 0.9323 - val_loss: 0.2903 - val_accuracy: 0.8841

Epoch 4/20

15000/15000 [=====] - 1s 45us/sample - loss: 0.1786 - accuracy: 0.9424 - val_loss: 0.2878 - val_accuracy: 0.8842

Epoch 5/20

15000/15000 [=====] - 1s 46us/sample - loss: 0.1430 - accuracy: 0.9543 - val_loss: 0.2809 - val_accuracy: 0.8860

Epoch 6/20

15000/15000 [=====] - 1s 45us/sample - loss: 0.1151 - accuracy: 0.9657 - val_loss: 0.2991 - val_accuracy: 0.8821

Epoch 7/20

15000/15000 [=====] - 1s 46us/sample - loss: 0.0961 - accuracy: 0.9714 - val_loss: 0.3104 - val_accuracy: 0.8822

Epoch 8/20

15000/15000 [=====] - 1s 47us/sample - loss: 0.0775 - accuracy: 0.9785 - val_loss: 0.3295 - val_accuracy: 0.8806

Epoch 9/20

15000/15000 [=====] - 1s 47us/sample - loss: 0.0670 - accuracy: 0.9819 - val_loss: 0.3484 - val_accuracy: 0.8812

Epoch 10/20

15000/15000 [=====] - 1s 46us/sample - loss: 0.0523 - accuracy: 0.9879 - val_loss: 0.3731 - val_accuracy: 0.8785

Epoch 11/20

15000/15000 [=====] - 1s 46us/sample - loss: 0.0407 - accuracy: 0.9910 - val_loss: 0.4024 - val_accuracy: 0.8742

Epoch 12/20

15000/15000 [=====] - 1s 46us/sample - loss: 0.0325 - accuracy: 0.9941 - val_loss: 0.4276 - val_accuracy: 0.8755

Epoch 13/20

15000/15000 [=====] - 1s 46us/sample - loss: 0.0298 - accuracy: 0.9928 - val_loss: 0.4601 - val_accuracy: 0.8717

Epoch 14/20

15000/15000 [=====] - 1s 46us/sample - loss: 0.0214 - accuracy: 0.9962 - val_loss: 0.4809 - val_accuracy: 0.8739

Epoch 15/20

15000/15000 [=====] - 1s 46us/sample - loss: 0.0133 - accuracy: 0.9989 - val_loss: 0.5154 - val_accuracy: 0.8707

Epoch 16/20

15000/15000 [=====] - 1s 46us/sample - loss: 0.0156 - accuracy: 0.9975 - val_loss: 0.5494 - val_accuracy: 0.8710

Epoch 17/20

15000/15000 [=====] - 1s 45us/sample - loss: 0.0103 - accuracy: 0.9990 - val_loss: 0.5808 - val_accuracy: 0.8681

Epoch 18/20

15000/15000 [=====] - 1s 46us/sample - loss: 0.0057 - accuracy: 0.9997 - val_loss: 0.6203 - val_accuracy: 0.8689

Epoch 19/20

15000/15000 [=====] - 1s 46us/sample - loss:

0.0085 - accuracy: 0.9985 - val_loss: 0.6468 - val_accuracy: 0.8672
 Epoch 20/20
 15000/15000 [=====] - 1s 43us/sample - loss:
 0.0032 - accuracy: 0.9999 - val_loss: 0.7158 - val_accuracy: 0.8603

```
In [11]: # Plotting the training and validation loss
import matplotlib.pyplot as plt

history_dict = history.history
loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']


epochs = range(1, len(loss_values) + 1)

plt.plot(epochs, loss_values, 'bo', label='Training loss')
plt.plot(epochs, val_loss_values, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```

<matplotlib.figure.Figure at 0x7fc22ef7e3c8>

```
In [12]: # Plotting the training and validation accuracy
plt.clf()
acc_values = history_dict['accuracy']
val_acc_values = history_dict['val_accuracy']
plt.plot(epochs, acc_values, 'bo', label='Training acc')
plt.plot(epochs, val_acc_values, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
In [17]: # Retraining a model from scratch
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(16, activation='relu', input_shape=(10000,)),
    tf.keras.layers.Dense(16, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')])

model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=[
    'accuracy'])

model.fit(partial_x_train, partial_y_train, epochs=4, batch_size=512, validation_data=(x_val, y_val))
results = model.evaluate(x_test, y_test)
```

Train on 15000 samples, validate on 10000 samples

Epoch 1/4

15000/15000 [=====] - 1s 84us/sample - loss: 0.5569 - accuracy: 0.7584 - val_loss: 0.4584 - val_accuracy: 0.8571

Epoch 2/4

15000/15000 [=====] - 1s 46us/sample - loss: 0.3633 - accuracy: 0.8943 - val_loss: 0.3461 - val_accuracy: 0.8806

Epoch 3/4

15000/15000 [=====] - 1s 46us/sample - loss: 0.2634 - accuracy: 0.9231 - val_loss: 0.3016 - val_accuracy: 0.8883

Epoch 4/4

15000/15000 [=====] - 1s 44us/sample - loss: 0.2076 - accuracy: 0.9382 - val_loss: 0.2808 - val_accuracy: 0.8900

25000/25000 [=====] - 2s 64us/sample - loss: 0.2976 - accuracy: 0.8832

```
In [18]: # Using a trained network to generate predictions on new data
model.predict(x_test)
```

```
Out[18]: array([[0.27379423],
                [0.9793424 ],
                [0.93266404],
                ...,
                [0.13444236],
                [0.17874026],
                [0.4376474 ]], dtype=float32)
```

In []: